

Егор Сеннов

Москва, Россия

☎ +7 (967)-174-69-30 | ✉ esennov2005@gmail.com | 💻 SennovE | 🌐 sennov

Summary

Backend developer with experience in content services. I write services and internal APIs in Go and Python, working with PostgreSQL, Redis, Kafka, and Docker/Kubernetes. I've participated in legacy service migrations, performance optimizations, and the launch of new microservices in production.

Experience

Т-Банк

Moscow, Russia

Backend Engineer

March 2025 - present

Media projects and content services team. Development and support of backend services and internal APIs.
Stack: Go, Python, PostgreSQL, Redis, Kafka, Docker, Kubernetes.

- Redesigned the backend airfare service for T-Journal after an MVP using NestJS: prepared RFCs, designed API contracts, a PostgreSQL schema, and integration with T-Travel, implemented the main backend logic in Go, and rolled it out to production. After the release, I was responsible for monitoring, analyzing integration incidents/errors, and further service improvements.
- I rewrote the legacy market data service for T-Journal using a new backend stack. I maintained compatibility with the current database schema and API so that the old and new implementations could be run side-by-side in production.
- Switched market data updates for securities, currencies, and other instruments to asynchronous processing. The ticker price update interval has been reduced from 3 hours to 15 minutes.
- Optimized the delivery of large lists of market data: reduced average response time and memory consumption by up to 10 times in certain product scenarios.
- I designed a survey response storage format based on JSON Schema generated from the form configuration. This allowed me to eliminate the need for separate backend structures and new API methods for each case, significantly reducing the time it took to create surveys.
- Refactored the legacy statistics service: switched synchronous database operations to async SQLAlchemy, updated logging and metrics, and removed over 60,000 lines of outdated handwritten and generated code. The update also reduced test run times by 15x.
- Replaced insecure serialization in the Redis cache with Pydantic/JSON. Prepared an ADR and migration plan, updated the code, tests, and Docker images.
- Participated in the development of open-source Python SDK for T-Investments : implemented a class generator from proto-files for typed gRPC requests.

Yandex

Moscow, Russia

Data Engineer

October 2024 - February 2025

The Plus and Fantech DWH content services team.
Stack: Python, SQL, ClickHouse.

- Improved ETL and analytical pipelines for content services: optimized SQL logic, reduced CPU consumption, and shortened the execution time of data processing chains.
- Participated in the implementation of ClickHouse into the data storage and aggregation pipeline. Migrated some product metric calculations from the previous data warehouse to ClickHouse.
- Developed loading of new data sources and transferred some calculations from backend services to the data platform.

Skills

- **Languages:** Go, Python, SQL, C++
- **Backend:** REST API, gRPC, microservices, API design, async processing, distributed systems
- **Datastores & Messaging:** PostgreSQL, MongoDB, Redis, Kafka, ClickHouse
- **Infrastructure:** Docker, Kubernetes, CI/CD, Linux
- **Practices:** service migration, performance optimization, observability, RFC/ADR

Projects

Examples of service architecture implementations in Go with different stacks.

- OpenAPI-first — A REST service in Go with code generation, PostgreSQL, JWT authentication, RBAC, CRUD for products, and business logic for orders/promo codes.
- REST + gRPC — A two-service booking system with gRPC interservice communication, separate databases, transactional seat reservations, Redis cache-aside, Redis Sentinel, retry/circuit breaker, and service-to-service authentication.
- Event Streaming + Analytics – an online cinema pipeline in Go with Kafka, Schema Registry, ClickHouse ingestion via Kafka Engine, PostgreSQL for ready-made metrics, event generation, DAU/retention/conversion aggregation, Grafana dashboard, export to S3/MinIO, and a fault-tolerant Kafka infrastructure.

Pre-commit hook for SQL code checking and command-line interface for formatting and receiving AI-based optimization recommendations.

- An AI-integrated command-line interface that optimizes SQL queries based on database analysis and resource usage. It generates a web page describing the improvements and highlighting query bottlenecks.
- A pre-commit hook that easily integrates into a project and automatically runs checks and optimizations on every commit.

Education

Bachelor’s degree in Applied Mathematics and Computer Science

Moscow, Russia

National Research University - Higher School of Economics

September 2023 - June 2027

- **Computer Science Courses:** Algorithms and data structures, operating systems, distributed systems, DBMS, cloud security, service-oriented architectures

School of Data Analysis

Moscow, Russia

Yandex

2025 - 2026

- Go

Backend Development School

Moscow, Russia

Yandex

Summer 2024

- Developed a backend platform for interactive longreads: implemented JWT authentication, longread versioning, file handling via S3, migrations, and automated tests for key functionality.
- Configured a fail-safe system with Nginx, load balancing, and automated deployment via CI/CD.

Extracurricular

Languages

- Russian (native)
- English (fluent)
- German (elementary)